

LESSON 23

MySQL & JPA Advanced

WEEK 05

Inheritance in JPA Entities

Introduction to Inheritance in JPA Entities

❖ Introduction to Inheritance in JPA Entities

❖ What is Inheritance in JPA?:

- JPA supports inheritance to model hierarchical entity classes, mapping Java OO inheritance to relational databases.
- Allows subclasses to inherit fields and relationships from a superclass.

❖ Inheritance Strategies:

- `@Inheritance(strategy = InheritanceType.SINGLE_TABLE)`: All classes in one table with discriminator.
- `@Inheritance(strategy = InheritanceType.TABLE_PER_CLASS)`: Separate table per concrete class.
- `@Inheritance(strategy = InheritanceType.JOINED)`: Separate table for superclass and each subclass (1:1 relationship via shared primary key).

❖ Focus on 1:1 Inheritance:

- Refers to JOINED strategy, where subclass tables link 1:1 to superclass table using shared PK.

❖ Reference: JPA Inheritance Specification

JOINED Inheritance Strategy (1:1 Mapping)

❖ **Definition:**

- Superclass has its own table; each subclass has a separate table with only subclass-specific fields.
- Subclass tables reference superclass table via shared primary key (1:1 relationship).

❖ **Annotations:**

- @Inheritance(strategy = InheritanceType.JOINED) on superclass.
- @PrimaryKeyJoinColumn optional for customizing join column.

❖ **Database Structure:**

- Superclass table: Common fields + PK.
- Subclass table: Subclass fields + PK (foreign key to superclass PK).

❖ **Explanation:**

- Queries join tables as needed; supports polymorphism (e.g., querying superclass returns mixed subclass instances).

Example of 1:1 Inheritance in JPA

❖ Generated Tables:

- person: id (PK), name.
- student: id (PK/FK to person.id), major.
- teacher: id (PK/FK to person.id), subject.

❖ Explanation:

- Inserting a Student creates rows in both person and student tables with same id.
- Query: SELECT p FROM Person p joins tables to fetch mixed Student/Teacher instances.

```
@Entity
@Inheritance(strategy = InheritanceType.JOINED)
public abstract class Person {
    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;
    private String name;
    // Getters and setters
}

@Entity
public class Student extends Person {
    private String major;
    // Getters and setters
}

@Entity
public class Teacher extends Person {
    private String subject;
    // Getters and setters
}
```

Advantages and Disadvantages of 1:1 Inheritance

❖ **Advantages:**

- Normalized database: No redundant fields; easy to add new subclasses.
- Supports polymorphism: Queries on superclass return subclass instances.
- Efficient for reads on specific subclasses (no unnecessary joins).

❖ **Disadvantages:**

- Performance overhead: Joins required for superclass queries.
- Complex inserts/updates: Multiple tables involved.
- Not suitable for deep hierarchies due to join complexity.

❖ **When to Use:**

- When normalization is important and hierarchies are not too deep.

Best Practices for 1:1 Inheritance in JPA

❖ Key Practices:

- Use @DiscriminatorColumn if needed for explicit type discrimination (though optional in JOINED).
- Optimize queries with Fetch Joins to avoid N+1 issues.
- Index join columns for better performance.
- Test polymorphism: Ensure repositories handle superclass queries correctly.

❖ Common Pitfalls:

- Overusing joins in deep hierarchies leading to slow queries.
- Forgetting to generate IDs in superclass.

```
@Query("SELECT p FROM Person p JOIN FETCH p WHERE p.id = :id")  
Person findByIdWithFetch(@Param("id") Long id);
```

JPA Entity Graphs for Dynamic Fetching

JPA Entity Graphs for Dynamic Fetching

❖ What is an Entity Graph?:

- Allows dynamic specification of which relationships to fetch (eager or lazy) for a query.
- Overrides default @FetchType settings in entity mappings.

❖ Types:

- @NamedEntityGraph: Defined statically in entity class.
- Dynamic Entity Graph: Created programmatically via EntityManager.

❖ Explanation:

- Fetches department eagerly for findAll(), overriding LAZY.
- Avoids N+1 issues by specifying related data in a single query.

JPA Entity Graphs for Dynamic Fetching

```
@Entity
@NamedEntityGraph(name = "Student.WithDepartmentAndCourses", attributeNodes = {
    @NamedAttributeNode("courses"),
    @NamedAttributeNode("department")
})
@NamedEntityGraph(name = "Student.WithCourses", attributeNodes = {
    @NamedAttributeNode("courses")
})
@NamedEntityGraph(name = "Student.WithDepartment", attributeNodes = {
    @NamedAttributeNode("department")
})
@Table(name = "students")
@FilterDef(name = "AvailableStudents", parameters = @ParamDef(name = "deleted", type = Boolean.class))
@Filter(name = "AvailableStudents", condition = "deleted = :deleted")
public class Student {
```

JPA Entity Graphs for Dynamic Fetching

```
// Avoid query N + 1 problem by using LEFT JOIN FETCH
```

```
@Query("SELECT s FROM Student s LEFT JOIN FETCH s.department LEFT JOIN FETCH s.courses")  
List<Student> getAllStudents();
```

```
// Avoid query N + 1 problem by using EntityGraph
```

```
@EntityGraph(attributePaths = { "department", "courses" })  
List<Student> findByStatus(StudentStatus status);
```

```
// Avoid query N + 1 problem by using EntityGraph
```

```
@EntityGraph(value = "Student.WithDepartmentAndCourses", type = EntityGraph.EntityGraphType.LOAD)  
List<Student> findByDeleted(boolean deleted);
```

JPA Entity Graphs for Dynamic Fetching

```
@PersistenceContext
private EntityManager em;

public List<StudentResponseDto> findByDepartmentId(Long departmentId) {
    // Using EntityManager to fetch students by department ID
    EntityGraph<?> graph = em.createEntityGraph(Student.class);
    graph.addAttributeNodes("department", "courses");
    List<Student> students = em
        .createQuery("SELECT s FROM Student s WHERE s.department.id = :departmentId", Student.class)
        .setHint("javax.persistence.loadgraph", graph) // Sử dụng Entity Graph
        .setParameter("departmentId", departmentId)
        .getResultList();

    // Convert to DTOs
    return students.stream()
        .map(this::convertToDto)
        .collect(Collectors.toList());
}
```

JPA Projections for Efficient Data Retrieval

JPA Projections for Efficient Data Retrieval

❖ What are Projections?:

- Retrieve only specific fields from entities instead of the entire object.
- Implemented via interfaces or DTO classes to reduce data transfer.

❖ Types:

- Interface-based: Define an interface with getter methods for desired fields.
- DTO-based: Use a custom class for complex projections.

❖ Explanation:

- Returns only id, name, and email, reducing data overhead.
- Works with Spring Data JPA's query derivation or custom @Query.

JPA Projections for Efficient Data Retrieval

```
public interface StudentProjection {  
    Long getId();  
    String getName();  
    String getEmail();  
}  
  
@Repository  
public interface StudentRepository extends JpaRepository<Student, Long> {  
    List<StudentProjection> findByNameContaining(String name);  
}
```

JPA Query Caching with Second-Level Cache

JPA Query Caching with Second-Level Cache

❖ What is Query Caching?:

- Caches query results to avoid repeated database hits for frequently executed queries.
- Uses Hibernate's second-level cache (L2 cache) with providers like EhCache.

❖ Setup:

- Enable L2 cache in application.properties:
 - `spring.jpa.properties.hibernate.cache.use_second_level_cache=true`
 - `spring.jpa.properties.hibernate.cache.region.factory_class=org.hibernate.cache.ehcache.EhCacheRegionFactory`

❖ Explanation:

- Caches query results and entities; subsequent calls retrieve from cache.
- Requires cache provider (e.g., EhCache) dependency in pom.xml.

JPA Query Caching with Second-Level Cache

❖ Add @Cacheable to entities:

```
@Entity
@Cacheable
@org.hibernate.annotations.Cache(usage = CacheConcurrencyStrategy.READ_WRITE)
public class Student {
    @Id
    @GeneratedValue
    private Long id;
    private String name;
}
```

JPA Query Caching with Second-Level Cache

❖ Query Cache Example:

```
@Repository
public interface StudentRepository extends JpaRepository<Student, Long> {
    @QueryHints(@QueryHint(name = "org.hibernate.cacheable", value = "true"))
    List<Student> findByName(String name);
}
```

JPA Lifecycle Events and Listeners

JPA Lifecycle Events and Listeners

❖ What are Lifecycle Events?:

- JPA triggers events during entity lifecycle: creation, update, deletion, etc.
- Annotations: @PrePersist, @PreUpdate, @PostLoad, @PostRemove, etc.

❖ Use Case:

- Automatically set creation/modification timestamps or enforce business rules.

❖ Explanation:

- ❖ @PrePersist: Executes before saving a new entity.
- ❖ @PreUpdate: Executes before updating an existing entity.

```
@Entity
public class Student {
    @Id
    @GeneratedValue
    private Long id;
    private String name;
    private LocalDateTime createdAt;
    private LocalDateTime updatedAt;

    @PrePersist
    public void prePersist() {
        createdAt = LocalDateTime.now();
        updatedAt = LocalDateTime.now();
    }

    @PreUpdate
    public void preUpdate() {
        updatedAt = LocalDateTime.now();
    }
}
```

JPA Native Queries

JPA Native Queries

❖ What are Native Queries?:

- Execute raw SQL queries when JPQL is insufficient for complex operations.
- Map results to entities, DTOs, or scalar values.

❖ Types:

- Entity-mapped: Return managed entities.
- Scalar/DTO-mapped: Return custom objects or raw data.

❖ Explanation:

- `nativeQuery = true`: Indicates raw SQL instead of JPQL.
- Useful for database-specific features or complex joins.

```
@Repository
public interface StudentRepository extends JpaRepository<Student, Long> {
    @Query(value = "...", nativeQuery = true)
    List<Object[]> findStudentAndDepartment(@Param("name") String name);
}
```

JPA Specifications for Dynamic Queries

JPA Specifications for Dynamic Queries

❖ What are Specifications?:

- Spring Data JPA feature to build dynamic, reusable query criteria.
- Uses Specification interface to define predicates for filtering.

❖ Use Case:

- Implement flexible search filters (e.g., search students by name, email, or department).

❖ Explanation:

- JpaSpecificationExecutor: Enables repository to use Specification.
- Combine multiple specifications with `and()`, `or()` for complex filters.

JPA Specifications for Dynamic Queries

```
import com.example.demo.entities.Student;
import org.springframework.data.jpa.domain.Specification;

public class StudentSpecifications {
    public static Specification<Student> hasName(String name) {
        return (root, query, cb) → cb.like(cb.lower(root.get("name")), "%" + name.toLowerCase() + "%");
    }
}
```

JPA Specifications for Dynamic Queries

```
import com.example.demo.entities.Student;
import org.springframework.data.jpa.domain.Specification;

public class StudentSpecifications {
    public static Specification<Student> hasName(String name) {
        return (root, query, cb) → cb.like(cb.lower(root.get("name")), "%" + name.toLowerCase() + "%");
    }
}

// JPA Specifications for Dynamic Queries
public List<Student> findByName(String name) {
    return this.studentJpaRepository.findAll(StudentSpecifications.hasName(name));
}
```

JPA Bulk Operations

Introduction to JPA Bulk Operations

❖ **Implementation:**

- Use @Query with @Modifying in repository methods.
- JPQL for entity-based updates; native SQL for complex cases.

❖ **Explanation:**

- @Modifying: Indicates the query modifies data (UPDATE/DELETE).
- Returns number of affected rows.
- Use @Transactional in service layer to ensure atomicity.

Introduction to JPA Bulk Operations

```
1  @Repository
2  public interface StudentJpaRepository extends JpaRepository<Student, Long>, JpaSpecificationExecutor<Student> {
3      @Query("SELECT s FROM Student s LEFT JOIN FETCH s.department")
4      List<Student> getAllStudentsWithDepartment();
5
6      @Modifying(clearAutomatically = true)
7      @Query("UPDATE Student s SET s.status = :status WHERE s.department.id = :departmentId")
8      int updateStudentStatus(@Param("status") String status, @Param("departmentId") Long departmentId);
9
10     @Modifying
11     @Query("DELETE FROM Student s WHERE s.status = :status")
12     int deleteInactiveStudents(@Param("status") String status);
13 }
14
```

Introduction to JPA Bulk Operations

@Repository

```
public interface StudentJpaRepository extends JpaRepository<Student, Long>, JpaSpecificationExecutor<Student> {  
    @Query("SELECT s FROM Student s LEFT JOIN FETCH s.department")  
    List<Student> getAllStudentsWithDepartment();  
  
    @Modifying(clearAutomatically = true)  
    @Query("UPDATE Student s SET s.status = :status WHERE s.department.id = :departmentId")  
    int updateStudentStatus(@Param("status") String status, @Param("departmentId") Long departmentId);  
  
    @Modifying  
    @Query("DELETE FROM Student s WHERE s.status = :status")  
    int deleteInactiveStudents(@Param("status") String status);  
}
```

@Transactional

```
public int updateStudentStatus(Long deptId, String status) {  
    return this.studentJpaRepository.updateStudentStatus(status, deptId);  
}
```

Soft Delete in JPA

Introduction to Soft Delete in JPA

❖ What is Soft Delete?:

- Marks records as "deleted" (e.g., via a flag) instead of physically removing them.
- Preserves data for auditing, recovery, or compliance.

❖ Why Use Soft Delete with Filters?:

- Automatically excludes "deleted" records from queries without modifying each query.
- Implemented using Hibernate filters in JPA.

❖ Key Components:

- A deleted flag in the entity (e.g., boolean or timestamp).
- @Filter and @FilterDef annotations for global filtering.

❖ Benefits:

- Simplifies queries; maintains data integrity.

Implementing Soft Delete with Filters

❖ Steps:

- Add a deleted field to the entity.
- Define @FilterDef and @Filter on the entity.
- Enable the filter in queries via EntityManager.

❖ Explanation:

- @FilterDef: Defines the filter with parameters.
- @Filter: Applies the condition (e.g., deleted = false).
- Enable filter per session for "active" records.

Implementing Soft Delete with Filters

```
@Transactional
public void softDeleteStudent(Long id) {
    Student student = this.studentJpaRepository.findById(id).orElseThrow();
    student.setDeleted(true);
    this.studentJpaRepository.save(student);
}
```

Implementing Soft Delete with Filters

```
@FilterDef(name = "softDeleteFilter", parameters = @ParamDef(name = "deleted", type = Boolean.class))
@Filter(name = "softDeleteFilter", condition = "deleted = :deleted")
public class Student {
    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;

    public List<StudentResponseDto> findAvailableStudents() {
        Session session = em.unwrap(Session.class);
        session.enableFilter("softDeleteFilter").setParameter("deleted", false);
        List<Student> student = em
            .createQuery("SELECT s FROM Student s LEFT JOIN FETCH s.department LEFT JOIN FETCH s.courses",
                Student.class)
            .getResultList();
        session.disableFilter("softDeleteFilter");
    }
}
```

Soft Delete Operations and Best Practices

❖ **Soft Delete Operation:**

- Set `deleted = true` instead of `repository.delete()`.

❖ **Restoring Records:**

- Set `deleted = false` to "undelete".

❖ **Best Practices:**

- Use a timestamp for `deletedAt` for better auditing.
- Enable filter globally via interceptor or aspect for consistency.
- Combine with projections to exclude deleted field in responses.

❖ **Pitfalls:**

- ❖ Forgotten filter enablement leads to including "deleted" records.
- ❖ Performance impact on large datasets; index the deleted column.

JPA Custom Converters

Introduction to JPA Custom Converters

❖ What are Custom Converters?:

- Map non-standard Java types to database columns.
- Use @Converter to define custom conversion logic.

❖ Use Cases:

- Enums to strings, JSON to text, custom objects to blobs.
- Ensures type safety and portability.

❖ Types:

- AttributeConverter: Implements AttributeConverter<X, Y> for entity attributes.
- Auto-apply or explicit via @Convert.

Implementing JPA Custom Converters

❖ Create enum

```
@Getter
public enum StudentStatus {
    ACTIVE(code:"ACT"), INACTIVE(code:"INA"), SUSPENDED(code:"SUS");
    private final String code;
    StudentStatus(String code) { this.code = code; }

    public static StudentStatus fromCode(String code) {
        for (StudentStatus s : StudentStatus.values()) {
            if (s.code.equals(code)) return s;
        }
        throw new IllegalArgumentException("Invalid code: " + code);
    }
}
```


Implementing JPA Custom Converters

❖ Create Converter

```
// autoApply = true: This converter will be applied to all fields of type StudentStatus
// autoApply = false: This converter must be explicitly specified in the entity field
// If you want to use this converter for a specific field,
// you must annotate that field with @Convert(converter = StudentStatusConverter.class)
@Converter(autoApply = false)
public class StudentStatusConverter implements AttributeConverter<StudentStatus, String> {
    @Override
    public String convertToDatabaseColumn(StudentStatus status) {
        return status != null ? status.getCode() : null;
    }

    @Override
    public StudentStatus convertToEntityAttribute(String code) {
        return code != null ? StudentStatus.fromCode(code) : null;
    }
}
```

Implementing JPA Custom Converters

❖ Apply to entity

```
public class Student {  
    @Id  
    @GeneratedValue(strategy = GenerationType.IDENTITY)  
    private Long id;  
    private String name;  
    private String email;  
    private String address;  
    private String password;  
  
    @Column  
    @Convert(converter = StudentStatusConverter.class)  
    private StudentStatus status;  
}
```

Conclusion and Next Steps

❖ **Summary:**

- Learned to build a CRUD application with Spring Boot, JPA, and MySQL.
- Covered entities, repositories, REST APIs, and basics JPA features.

❖ **Next Steps:**

- Build a full-stack application with a front-end (e.g., React).

❖ **References:**

- Spring Boot
- Spring Data JPA
- JPA Specification