

LESSON 23

MySQL & JPA Basics

WEEK 05

What is JPA?

❖ **Definition:**

- Java Persistence API (JPA) is a specification for object-relational mapping (ORM).
- Maps Java objects to database tables.

❖ **Key Components:**

- Entity, EntityManager, Persistence Unit.

❖ **Why Use JPA?:**

- Simplifies database operations (no manual SQL for basic CRUD).
- Supports multiple databases (e.g., MySQL, PostgreSQL).

Spring Data JPA Overview

❖ What is Spring Data JPA?:

- Extension of Spring Data for JPA-based repositories.
- Provides built-in methods for CRUD operations.

❖ Key Features:

- Repository interfaces (CrudRepository, JpaRepository).
- Query methods derived from method names.
- Custom queries with @Query annotation.

❖ Benefits:

- Reduces boilerplate code for database access.
- Integrates seamlessly with Spring Boot.

Setting Up the Development Environment

❖ **Tools Required:**

- JDK 17+, IntelliJ IDEA (or Eclipse or VS Code), MySQL, Maven/Gradle.

❖ **Steps:**

- Install MySQL and create a database (e.g., school_db).
- Configure IDE with Spring Boot plugin.
- Add Spring Boot Starter dependencies.

❖ **Dependencies:**

- spring-boot-starter-data-jpa
- mysql-connector-java

❖ **Reference:** Spring Initializr

Creating a Spring Boot Project

❖ Using Spring Initializr:

- Select Java, Gradle, Spring Boot 3.x.
- Add dependencies: Spring Web, Spring Data JPA, MySQL Driver.

❖ Project Structure:

- src/main/java: Application code.
- src/main/resources: Configuration files (e.g., application.properties).

❖ Example:

- Generate project at start.spring.io.
- Import into IDE and run.

Configuring MySQL in Spring Boot

❖ **Configuration File:**

- Edit application.properties to connect to MySQL.

❖ **Example Configuration:**

```
spring.datasource.url=jdbc:mysql://localhost:3306/school_db  
spring.datasource.username=root  
spring.datasource.password=your_password  
spring.jpa.hibernate.ddl-auto=update
```

❖ **Explanation:**

- ddl-auto=update: Automatically creates/updates database schema based on entities.
- Ensure MySQL server is running.

Creating a JPA Entity

❖ What is an Entity?:

- A Java class mapped to a database table.

❖ Annotations:

- @Entity: Marks class as an entity.
- @Id: Defines primary key.
- @GeneratedValue: Auto-generates ID values.

Creating a JPA Repository

❖ **Repository Interface:**

- Extends `JpaRepository<EntityClass, IDType>`.
- Provides built-in CRUD methods.

❖ **Example:**

```
public interface StudentRepository extends JpaRepository<Student, Long> {  
    // Custom query methods  
}
```

❖ **Built-in Methods:**

`save()`, `findById()`, `findAll()`, `deleteById()`.

Implementing Create Operation

❖ Purpose:

- Save a new entity to the database.

❖ Explanation:

- save() persists the entity to the database.
- Returns the saved entity with generated ID.

```
@Autowired  
private StudentRepository repository;  
  
public Student createStudent(Student student) {  
    return repository.save(student);  
}
```

Implementing Read Operation

❖ Purpose:

- Retrieve entities from the database.

❖ Explanation:

- findAll(): Retrieves all records.
- findById(): Retrieves a single record by ID.

```
public List<Student> getAllStudents() {  
    return repository.findAll();  
}
```

```
public Optional<Student> getStudentById(Long id) {  
    return repository.findById(id);  
}
```

Implementing Update Operation

❖ Purpose:

- Modify an existing entity in the database.

❖ Explanation:

- Fetch entity, update fields, and save.

```
public Student updateStudent(Long id, Student updatedStudent) {  
    Student student = repository.findById(id).orElseThrow();  
    student.setName(updatedStudent.getName());  
    student.setEmail(updatedStudent.getEmail());  
    return repository.save(student);  
}
```

Implementing Delete Operation

❖ Purpose:

- Remove an entity from the database.

❖ Explanation:

- deleteById() removes the entity with the specified ID.
- Throws exception if ID does not exist.

```
public void deleteStudent(Long id) {  
    repository.deleteById(id);  
}
```

Creating a REST Controller

❖ Purpose:

- Expose CRUD operations via RESTful APIs.

```
@RestController
@RequestMapping("/api/students")
public class StudentController {
    @Autowired
    private StudentService service;

    @PostMapping
    public Student create(@RequestBody Student student) {
        return service.createStudent(student);
    }

    @GetMapping
    public List<Student> getAll() {
        return service.getAllStudents();
    }
}
```

Full CRUD Example

❖ Scenario:

- Manage student records (create, read, update, delete).

```
@RestController
@RequestMapping("/api/students")
public class StudentController {

    @Autowired
    private StudentService service;

    @PostMapping
    public Student create(@RequestBody Student student) {
        return service.createStudent(student);
    }

    @GetMapping("/{id}")
    public Student getById(@PathVariable Long id) {
        return service.getStudentById(id).orElseThrow();
    }

    @PutMapping("/{id}")
    public Student update(@PathVariable Long id, @RequestBody Student student) {
        return service.updateStudent(id, student);
    }

    @DeleteMapping("/{id}")
    public void delete(@PathVariable Long id) {
        service.deleteStudent(id);
    }
}
```

Testing APIs with Postman

❖ Steps:

- Start Spring Boot application.
- Use Postman to send HTTP requests (POST, GET, PUT, DELETE).

❖ Example:

- POST: `http://localhost:8080/api/students` with JSON body:

```
{"name": "John Doe", "email": "john@example.com"}
```

Validation with JPA

❖ Purpose:

- Ensure valid data before saving to database.

❖ Explanation:

- Use annotations like @NotNull, @Email from javax.validation.

```
public class Student {  
    @NotNull  
    private String name;  
  
    @Email  
    private String email;  
  
}
```


Introduction to JPA Relationship Annotations

❖ Purpose:

- Define relationships between entities (e.g., Student, Department) in JPA.
- Annotations: @OneToOne, @OneToMany, @ManyToOne, @ManyToMany.

❖ Key Concepts:

- Owning side: Defines the relationship (owns the foreign key).
- Inverse side: References the owning side (uses mappedBy).
- Cascade and fetch strategies control behavior and performance.

❖ Why Important?:

- Enables modeling of complex data relationships in the database.
- Simplifies querying and data management.

Configuring @OneToOne Relationship

❖ Definition:

- One entity instance is related to exactly one instance of another entity.
- Example: A Student has one Address.

❖ Explanation:

- @JoinColumn: Specifies the foreign key column (address_id) in the Student table.
- cascade: Propagates operations (e.g., save, delete) to the related entity.

```
@Entity
public class Student {
    @Id
    @GeneratedValue
    private Long id;
    private String name;

    @OneToOne(cascade = CascadeType.ALL)
    @JoinColumn(name = "address_id")
    private Address address;
}

@Entity
public class Address {
    @Id
    @GeneratedValue
    private Long id;
    private String street;
}
```

Configuring @OneToMany and @ManyToOne

❖ Definition:

- @OneToMany: One entity relates to multiple instances of another (e.g., one Department has many Students).
- @ManyToOne: Many instances relate to one instance (e.g., many Students belong to one Department).

❖ Explanation:

- @ManyToOne (owning side): Defines the foreign key (department_id) in the Student table.
- @OneToMany (inverse side): Uses mappedBy to reference the owning side.

```
@Entity
public class Department {
    @Id
    @GeneratedValue
    private Long id;
    private String name;

    @OneToMany(mappedBy = "department", cascade = CascadeType.ALL)
    private List<Student> students = new ArrayList<>();
}

@Entity
public class Student {
    @Id
    @GeneratedValue
    private Long id;
    private String name;

    @ManyToOne
    @JoinColumn(name = "department_id")
    private Department department;
}
```

Configuring @ManyToMany Relationship

❖ Definition:

- Multiple instances of one entity relate to multiple instances of another.
- Example: Students enroll in multiple Courses, and Courses have multiple Students.

❖ Explanation:

- @JoinTable: Defines the join table (student_course) with foreign keys.
- mappedBy: Specifies the owning side (Student) to avoid duplicate mappings.

```
@Entity
public class Student {
    @Id
    @GeneratedValue
    private Long id;
    private String name;

    @ManyToMany(cascade = CascadeType.ALL)
    @JoinTable(
        name = "student_course",
        joinColumns = @JoinColumn(name = "student_id"),
        inverseJoinColumns = @JoinColumn(name = "course_id")
    )
    private List<Course> courses = new ArrayList<>();
}

@Entity
public class Course {
    @Id
    @GeneratedValue
    private Long id;
    private String title;

    @ManyToMany(mappedBy = "courses")
    private List<Student> students = new ArrayList<>();
}
```

Cascade and Fetch Strategies in Relationships

❖ Cascade:

- Controls propagation of operations (e.g., save, delete) to related entities.
- Options: CascadeType.ALL, PERSIST, MERGE, REMOVE, etc.
- Example: cascade = CascadeType.ALL saves related entities automatically.

❖ Fetch Strategies:

- FetchType.LAZY: Loads related data only when accessed (default for @OneToMany, @ManyToMany).
- FetchType.EAGER: Loads related data immediately (default for @ManyToOne, @OneToOne).

```
@OneToMany(mappedBy = "department", cascade = CascadeType.PERSIST, fetch = FetchType.LAZY)  
private List<Student> students;
```

Introduction to Joins in JPA

❖ What are Joins in JPA?:

- Joins in JPA are used to combine data from multiple entities based on relationships.
- Defined in JPQL (Java Persistence Query Language) or Criteria API.
- Support for INNER JOIN, LEFT JOIN, RIGHT JOIN, and implicit joins.

❖ Why Use Joins?:

- Retrieve related data in a single query, avoiding multiple database calls.
- Essential for querying associations like @OneToMany or @ManyToMany.

Inner Join in JPA

❖ Definition:

- Returns records that have matching values in both entities.
- Equivalent to SQL INNER JOIN.

❖ Explanation:

- Joins Student and Department entities on their relationship.
- Only includes students with a matching department.

```
@Query("SELECT s FROM Student s JOIN s.department d WHERE d.name = :deptName")  
List<Student> findStudentsByDepartment(@Param("deptName") String deptName);
```

Left Outer Join in JPA

❖ Definition:

- Returns all records from the left entity and matching records from the right.
- Non-matching right records are null.

❖ Explanation:

- Includes all students, even those without a department.
- Useful for optional relationships.

```
@Query("SELECT s FROM Student s LEFT JOIN s.department d WHERE d.name = :deptName OR d IS NULL")  
List<Student> findStudentsWithOptionalDepartment(@Param("deptName") String deptName);
```


Right Outer Join in JPA

❖ Definition:

- Returns all records from the right entity and matching records from the left.
- Non-matching left records are null.

❖ Explanation:

- Includes all departments, even those without students.
- Less common than LEFT JOIN but symmetric.

```
@Query("SELECT d FROM Department d RIGHT JOIN d.students s WHERE s.name = :studentName")  
List<Department> findDepartmentsByStudent(@Param("studentName") String studentName);
```

Right Outer Join in JPA

❖ Definition:

- Returns all records from the right entity and matching records from the left.
- Non-matching left records are null.

```
@Query("SELECT d FROM Department d RIGHT JOIN d.students s WHERE s.name = :studentName")  
List<Department> findDepartmentsByStudent(@Param("studentName") String studentName);
```

Fetch Joins in JPA

❖ Definition:

- Eagerly fetches related entities in a single query to avoid N+1 problem.
- Uses FETCH keyword in JPQL.

❖ Explanation:

- Loads departments immediately, preventing lazy loading exceptions.
- Improves performance for read operations.

```
@Query("SELECT s FROM Student s JOIN FETCH s.department d")  
List<Student> findAllStudentsWithDepartments();
```

Introduction to Paging in JPA

❖ Introduction to Paging in JPA

❖ What is Paging?:

- Paging allows retrieving large datasets in smaller chunks (pages) to improve performance and usability.
- Essential for applications with large databases to avoid loading all data at once.

❖ Key Components in Spring Data JPA:

- Pageable: Interface for pagination and sorting information.
- Page: Represents a page of data with metadata (total pages, total elements).
- Slice: Similar to Page but without total count (faster for large datasets).

❖ Why Use Paging?:

- Reduces memory usage, improves response times, and enables features like infinite scrolling.

Using Pageable in JpaRepository

❖ Pageable Interface:

- Created using `PageRequest.of(pageNumber, pageSize, sort)` to specify page index, size, and sorting.

❖ Repository Methods:

- Extend `JpaRepository` and add methods returning `Page<T>` or `Slice<T>`.

```
public interface StudentRepository extends JpaRepository<Student, Long> {  
    Page<Student> findAll(Pageable pageable);  
}
```

```
Pageable pageable = PageRequest.of(0, 10, Sort.by("name").ascending());  
Page<Student> studentsPage = repository.findAll(pageable);
```

Working with Page and Slice

❖ Page:

- Provides full pagination info: content, total pages, total elements.
- Example: `studentsPage.getTotalElements()`, `studentsPage.getTotalPages()`.

❖ Slice:

- Lighter than Page; no total count (avoids expensive COUNT queries).
- Use for "load more" features where total is not needed.

```
Slice<Student> studentsSlice = repository.findAll(Pageable pageable);  
List<Student> content = studentsSlice.getContent();  
boolean hasNext = studentsSlice.hasNext();
```

Custom Queries with Paging

❖ Custom JPQL Queries:

- Use @Query with Pageable for custom pagination.

❖ Explanation:

- Pageable is appended as the last parameter in custom queries.
- Supports sorting and pagination on derived or custom queries.

```
public interface StudentRepository extends JpaRepository<Student, Long> {  
    @Query("SELECT s FROM Student s WHERE s.name LIKE %:name%")  
    Page<Student> findByNameContaining(@Param("name") String name, Pageable pageable);  
}
```

```
Page<Student> results = repository.findByNameContaining("John", pageable);
```

Best Practices for Paging in JPA

❖ Key Practices:

- Use Slice for large datasets to avoid slow COUNT queries.
- Combine with sorting: `Sort.by("field").ascending()` for user-friendly results.
- Handle edge cases: Empty pages, invalid page numbers.
- Use in REST APIs: Return Page metadata in responses for client-side pagination.

❖ Common Pitfalls:

- N+1 queries: Use Fetch Joins with paging for relationships.
- Performance: Index columns used in sorting/filters.

LESSON 23

MySQL & JPA Advanced

WEEK 05

Introduction to Inheritance in JPA Entities

❖ Introduction to Inheritance in JPA Entities

❖ What is Inheritance in JPA?:

- JPA supports inheritance to model hierarchical entity classes, mapping Java OO inheritance to relational databases.
- Allows subclasses to inherit fields and relationships from a superclass.

❖ Inheritance Strategies:

- `@Inheritance(strategy = InheritanceType.SINGLE_TABLE)`: All classes in one table with discriminator.
- `@Inheritance(strategy = InheritanceType.TABLE_PER_CLASS)`: Separate table per concrete class.
- `@Inheritance(strategy = InheritanceType.JOINED)`: Separate table for superclass and each subclass (1:1 relationship via shared primary key).

❖ Focus on 1:1 Inheritance:

- Refers to JOINED strategy, where subclass tables link 1:1 to superclass table using shared PK.

❖ Reference: JPA Inheritance Specification

JOINED Inheritance Strategy (1:1 Mapping)

❖ **Definition:**

- Superclass has its own table; each subclass has a separate table with only subclass-specific fields.
- Subclass tables reference superclass table via shared primary key (1:1 relationship).

❖ **Annotations:**

- @Inheritance(strategy = InheritanceType.JOINED) on superclass.
- @PrimaryKeyJoinColumn optional for customizing join column.

❖ **Database Structure:**

- Superclass table: Common fields + PK.
- Subclass table: Subclass fields + PK (foreign key to superclass PK).

❖ **Explanation:**

- Queries join tables as needed; supports polymorphism (e.g., querying superclass returns mixed subclass instances).

Example of 1:1 Inheritance in JPA

❖ Generated Tables:

- person: id (PK), name.
- student: id (PK/FK to person.id), major.
- teacher: id (PK/FK to person.id), subject.

❖ Explanation:

- Inserting a Student creates rows in both person and student tables with same id.
- Query: SELECT p FROM Person p joins tables to fetch mixed Student/Teacher instances.

❖ Reference: Thorben Janssen: JPA Joined Strategy

```
@Entity
@Inheritance(strategy = InheritanceType.JOINED)
public abstract class Person {
    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;
    private String name;
    // Getters and setters
}

@Entity
public class Student extends Person {
    private String major;
    // Getters and setters
}

@Entity
public class Teacher extends Person {
    private String subject;
    // Getters and setters
}
```

Advantages and Disadvantages of 1:1 Inheritance

❖ **Advantages:**

- Normalized database: No redundant fields; easy to add new subclasses.
- Supports polymorphism: Queries on superclass return subclass instances.
- Efficient for reads on specific subclasses (no unnecessary joins).

❖ **Disadvantages:**

- Performance overhead: Joins required for superclass queries.
- Complex inserts/updates: Multiple tables involved.
- Not suitable for deep hierarchies due to join complexity.

❖ **When to Use:**

- When normalization is important and hierarchies are not too deep.

Best Practices for 1:1 Inheritance in JPA

❖ Key Practices:

- Use @DiscriminatorColumn if needed for explicit type discrimination (though optional in JOINED).
- Optimize queries with Fetch Joins to avoid N+1 issues.
- Index join columns for better performance.
- Test polymorphism: Ensure repositories handle superclass queries correctly.

❖ Common Pitfalls:

- Overusing joins in deep hierarchies leading to slow queries.
- Forgetting to generate IDs in superclass.

```
@Query("SELECT p FROM Person p JOIN FETCH p WHERE p.id = :id")  
Person findByIdWithFetch(@Param("id") Long id);
```

JPA Entity Graphs for Dynamic Fetching

❖ What is an Entity Graph?:

- Allows dynamic specification of which relationships to fetch (eager or lazy) for a query.
- Overrides default @FetchType settings in entity mappings.

❖ Types:

- @NamedEntityGraph: Defined statically in entity class.
- Dynamic Entity Graph: Created programmatically via EntityManager.

❖ Explanation:

- Fetches department eagerly for findAll(), overriding LAZY.
- Avoids N+1 issues by specifying related data in a single query.

JPA Entity Graphs for Dynamic Fetching

```
@Entity
@NamedEntityGraph(name = "Student.withDepartment",
                    attributeNodes = @NamedAttributeNode("department"))
public class Student {
    @Id
    @GeneratedValue
    private Long id;
    private String name;
    @ManyToOne(fetch = FetchType.LAZY)
    private Department department;
}

@Repository
public interface StudentRepository extends JpaRepository<Student, Long> {
    @EntityGraph(value = "Student.withDepartment")
    List<Student> findAll();
}
```


JPA Projections for Efficient Data Retrieval

❖ What are Projections?:

- Retrieve only specific fields from entities instead of the entire object.
- Implemented via interfaces or DTO classes to reduce data transfer.

❖ Types:

- Interface-based: Define an interface with getter methods for desired fields.
- DTO-based: Use a custom class for complex projections.

❖ Explanation:

- Returns only id, name, and email, reducing data overhead.
- Works with Spring Data JPA's query derivation or custom @Query.

JPA Projections for Efficient Data Retrieval

```
public interface StudentProjection {  
    Long getId();  
    String getName();  
    String getEmail();  
}  
  
@Repository  
public interface StudentRepository extends JpaRepository<Student, Long> {  
    List<StudentProjection> findByNameContaining(String name);  
}
```

JPA Query Caching with Second-Level Cache

❖ What is Query Caching?:

- Caches query results to avoid repeated database hits for frequently executed queries.
- Uses Hibernate's second-level cache (L2 cache) with providers like EhCache.

❖ Setup:

- Enable L2 cache in application.properties:
 - `spring.jpa.properties.hibernate.cache.use_second_level_cache=true`
 - `spring.jpa.properties.hibernate.cache.region.factory_class=org.hibernate.cache.ehcache.EhCacheRegionFactory`

❖ Explanation:

- Caches query results and entities; subsequent calls retrieve from cache.
- Requires cache provider (e.g., EhCache) dependency in pom.xml.

JPA Query Caching with Second-Level Cache

❖ Add @Cacheable to entities:

```
@Entity
@Cacheable
@org.hibernate.annotations.Cache(usage = CacheConcurrencyStrategy.READ_WRITE)
public class Student {
    @Id
    @GeneratedValue
    private Long id;
    private String name;
}
```

JPA Query Caching with Second-Level Cache

❖ Query Cache Example:

```
@Repository
public interface StudentRepository extends JpaRepository<Student, Long> {
    @QueryHints(@QueryHint(name = "org.hibernate.cacheable", value = "true"))
    List<Student> findByName(String name);
}
```

JPA Lifecycle Events and Listeners

❖ What are Lifecycle Events?:

- JPA triggers events during entity lifecycle: creation, update, deletion, etc.
- Annotations: @PrePersist, @PreUpdate, @PostLoad, @PostRemove, etc.

❖ Use Case:

- Automatically set creation/modification timestamps or enforce business rules.

❖ Explanation:

- ❖ @PrePersist: Executes before saving a new entity.
- ❖ @PreUpdate: Executes before updating an existing entity.

```
@Entity
public class Student {
    @Id
    @GeneratedValue
    private Long id;
    private String name;
    private LocalDateTime createdAt;
    private LocalDateTime updatedAt;

    @PrePersist
    public void prePersist() {
        createdAt = LocalDateTime.now();
        updatedAt = LocalDateTime.now();
    }

    @PreUpdate
    public void preUpdate() {
        updatedAt = LocalDateTime.now();
    }
}
```

JPA Native Queries

❖ What are Native Queries?:

- Execute raw SQL queries when JPQL is insufficient for complex operations.
- Map results to entities, DTOs, or scalar values.

❖ Types:

- Entity-mapped: Return managed entities.
- Scalar/DTO-mapped: Return custom objects or raw data.

❖ Explanation:

- `nativeQuery = true`: Indicates raw SQL instead of JPQL.
- Useful for database-specific features or complex joins.

```
@Repository
public interface StudentRepository extends JpaRepository<Student, Long> {
    @Query(value = "...", nativeQuery = true)
    List<Object[]> findStudentAndDepartment(@Param("name") String name);
}
```

JPA Specifications for Dynamic Queries

❖ What are Specifications?:

- Spring Data JPA feature to build dynamic, reusable query criteria.
- Uses Specification interface to define predicates for filtering.

❖ Use Case:

- Implement flexible search filters (e.g., search students by name, email, or department).

❖ Explanation:

- JpaSpecificationExecutor: Enables repository to use Specification.
- Combine multiple specifications with `and()`, `or()` for complex filters.

Introduction to JPA Bulk Operations

❖ **Implementation:**

- Use @Query with @Modifying in repository methods.
- JPQL for entity-based updates; native SQL for complex cases.

❖ **Explanation:**

- @Modifying: Indicates the query modifies data (UPDATE/DELETE).
- Returns number of affected rows.
- Use @Transactional in service layer to ensure atomicity.

Introduction to JPA Bulk Operations

```
1  @Repository
2  public interface StudentJpaRepository extends JpaRepository<Student, Long>, JpaSpecificationExecutor<Student> {
3      @Query("SELECT s FROM Student s LEFT JOIN FETCH s.department")
4      List<Student> getAllStudentsWithDepartment();
5
6      @Modifying(clearAutomatically = true)
7      @Query("UPDATE Student s SET s.status = :status WHERE s.department.id = :departmentId")
8      int updateStudentStatus(@Param("status") String status, @Param("departmentId") Long departmentId);
9
10     @Modifying
11     @Query("DELETE FROM Student s WHERE s.status = :status")
12     int deleteInactiveStudents(@Param("status") String status);
13 }
14
```

Introduction to JPA Bulk Operations

@Repository

```
public interface StudentJpaRepository extends JpaRepository<Student, Long>, JpaSpecificationExecutor<Student> {  
    @Query("SELECT s FROM Student s LEFT JOIN FETCH s.department")  
    List<Student> getAllStudentsWithDepartment();  
  
    @Modifying(clearAutomatically = true)  
    @Query("UPDATE Student s SET s.status = :status WHERE s.department.id = :departmentId")  
    int updateStudentStatus(@Param("status") String status, @Param("departmentId") Long departmentId);  
  
    @Modifying  
    @Query("DELETE FROM Student s WHERE s.status = :status")  
    int deleteInactiveStudents(@Param("status") String status);  
}
```

@Transactional

```
public int updateStudentStatus(Long deptId, String status) {  
    return this.studentJpaRepository.updateStudentStatus(status, deptId);  
}
```

Introduction to Soft Delete in JPA

❖ What is Soft Delete?:

- Marks records as "deleted" (e.g., via a flag) instead of physically removing them.
- Preserves data for auditing, recovery, or compliance.

❖ Why Use Soft Delete with Filters?:

- Automatically excludes "deleted" records from queries without modifying each query.
- Implemented using Hibernate filters in JPA.

❖ Key Components:

- A deleted flag in the entity (e.g., boolean or timestamp).
- @Filter and @FilterDef annotations for global filtering.

❖ Benefits:

- Simplifies queries; maintains data integrity.

Implementing Soft Delete with Filters

❖ Steps:

- Add a deleted field to the entity.
- Define @FilterDef and @Filter on the entity.
- Enable the filter in queries via EntityManager.

❖ Explanation:

- @FilterDef: Defines the filter with parameters.
- @Filter: Applies the condition (e.g., deleted = false).
- Enable filter per session for "active" records.

Implementing Soft Delete with Filters

```
@Entity
@FilterDef(name = "activeOnly", parameters = @ParamDef(name = "active", type = Boolean.class))
@Filter(name = "activeOnly", condition = "deleted = :active")
public class Student {
    ...
}

@Service
public class StudentService {
    @PersistenceContext
    private EntityManager em;

    public List<Student> findActiveStudents() {
        em.unwrap(Session.class).enableFilter("activeOnly").setParameter("active", false);
        return em.createQuery("FROM Student", Student.class).getResultList();
    }
}
```

Soft Delete Operations and Best Practices

❖ **Soft Delete Operation:**

- Set deleted = true instead of repository.delete().

❖ **Restoring Records:**

- Set deleted = false to "undelete".

❖ **Best Practices:**

- Use a timestamp for deletedAt for better auditing.
- Enable filter globally via interceptor or aspect for consistency.
- Combine with projections to exclude deleted field in responses.

❖ **Pitfalls:**

- ❖ Forgotten filter enablement leads to including "deleted" records.
- ❖ Performance impact on large datasets; index the deleted column.

```
@Transactional
```

```
public void softDeleteStudent(Long id) {  
    Student student = repository.findById(id).orElseThrow();  
    student.setDeleted(true);  
    repository.save(student);  
}
```

Introduction to JPA Custom Converters

❖ What are Custom Converters?:

- Map non-standard Java types to database columns.
- Use @Converter to define custom conversion logic.

❖ Use Cases:

- Enums to strings, JSON to text, custom objects to blobs.
- Ensures type safety and portability.

❖ Types:

- AttributeConverter: Implements AttributeConverter<X, Y> for entity attributes.
- Auto-apply or explicit via @Convert.

Implementing JPA Custom Converters

❖ Create enum

```
@Getter
public enum StudentStatus {
    ACTIVE(code:"ACT"), INACTIVE(code:"INA"), SUSPENDED(code:"SUS");
    private final String code;
    StudentStatus(String code) { this.code = code; }

    public static StudentStatus fromCode(String code) {
        for (StudentStatus s : StudentStatus.values()) {
            if (s.code.equals(code)) return s;
        }
        throw new IllegalArgumentException("Invalid code: " + code);
    }
}
```

Implementing JPA Custom Converters

❖ Create Converter

```
// autoApply = true: This converter will be applied to all fields of type StudentStatus
// autoApply = false: This converter must be explicitly specified in the entity field
// If you want to use this converter for a specific field,
// you must annotate that field with @Convert(converter = StudentStatusConverter.class)
@Converter(autoApply = false)
public class StudentStatusConverter implements AttributeConverter<StudentStatus, String> {
    @Override
    public String convertToDatabaseColumn(StudentStatus status) {
        return status != null ? status.getCode() : null;
    }

    @Override
    public StudentStatus convertToEntityAttribute(String code) {
        return code != null ? StudentStatus.fromCode(code) : null;
    }
}
```

Implementing JPA Custom Converters

❖ Apply to entity

```
public class Student {  
    @Id  
    @GeneratedValue(strategy = GenerationType.IDENTITY)  
    private Long id;  
    private String name;  
    private String email;  
    private String address;  
    private String password;  
  
    @Column  
    @Convert(converter = StudentStatusConverter.class)  
    private StudentStatus status;
```

Conclusion and Next Steps

❖ **Summary:**

- Learned to build a CRUD application with Spring Boot, JPA, and MySQL.
- Covered entities, repositories, REST APIs, and basics JPA features.

❖ **Next Steps:**

- Build a full-stack application with a front-end (e.g., React).

❖ **References:**

- Spring Boot
- Spring Data JPA
- JPA Specification