

# LESSON 23

## MySQL & JPA Basics

### WEEK 05

# What is JPA?

## ❖ **Definition:**

- Java Persistence API (JPA) is a specification for object-relational mapping (ORM).
- Maps Java objects to database tables.

## ❖ **Key Components:**

- Entity, EntityManager, Persistence Unit.

## ❖ **Why Use JPA?:**

- Simplifies database operations (no manual SQL for basic CRUD).
- Supports multiple databases (e.g., MySQL, PostgreSQL).

# Spring Data JPA Overview

## ❖ What is Spring Data JPA?:

- Extension of Spring Data for JPA-based repositories.
- Provides built-in methods for CRUD operations.

## ❖ Key Features:

- Repository interfaces (CrudRepository, JpaRepository).
- Query methods derived from method names.
- Custom queries with @Query annotation.

## ❖ Benefits:

- Reduces boilerplate code for database access.
- Integrates seamlessly with Spring Boot.

# Setting Up the Development Environment

## ❖ **Tools Required:**

- JDK 17+, IntelliJ IDEA (or Eclipse or VS Code), MySQL, Maven/Gradle.

## ❖ **Steps:**

- Install MySQL and create a database (e.g., school\_db).
- Configure IDE with Spring Boot plugin.
- Add Spring Boot Starter dependencies.

## ❖ **Dependencies:**

- spring-boot-starter-data-jpa
- mysql-connector-java

## ❖ **Reference:** Spring Initializr

# Creating a Spring Boot Project

## ❖ Using Spring Initializr:

- Select Java, Gradle, Spring Boot 3.x.
- Add dependencies: Spring Web, Spring Data JPA, MySQL Driver.

## ❖ Project Structure:

- src/main/java: Application code.
- src/main/resources: Configuration files (e.g., application.properties).

## ❖ Example:

- Generate project at [start.spring.io](https://start.spring.io).
- Import into IDE and run.

# Configuring MySQL in Spring Boot

## ❖ Configuration File:

- Edit application.properties to connect to MySQL.

## ❖ Example Configuration:

```
spring.datasource.url=jdbc:mysql://localhost:3306/school_db  
spring.datasource.username=root  
spring.datasource.password=your_password  
spring.jpa.hibernate.ddl-auto=update
```

## ❖ Explanation:

- ddl-auto=update: Automatically creates/updates database schema based on entities.
- Ensure MySQL server is running.

# Creating a JPA Entity

## ❖ What is an Entity?:

- A Java class mapped to a database table.

## ❖ Annotations:

- @Entity: Marks class as an entity.
- @Id: Defines primary key.
- @GeneratedValue: Auto-generates ID values.

# Creating a JPA Repository

## ❖ **Repository Interface:**

- Extends `JpaRepository<EntityClass, IDType>`.
- Provides built-in CRUD methods.

## ❖ **Example:**

```
public interface StudentRepository extends JpaRepository<Student, Long> {  
    // Custom query methods  
}
```

## ❖ **Built-in Methods:**

`save()`, `findById()`, `findAll()`, `deleteById()`.



# Implementing Create Operation

## ❖ Purpose:

- Save a new entity to the database.

## ❖ Explanation:

- save() persists the entity to the database.
- Returns the saved entity with generated ID.

```
@Autowired  
private StudentRepository repository;  
  
public Student createStudent(Student student) {  
    return repository.save(student);  
}
```

# Implementing Read Operation

## ❖ Purpose:

- Retrieve entities from the database.

## ❖ Explanation:

- findAll(): Retrieves all records.
- findById(): Retrieves a single record by ID.

```
public List<Student> getAllStudents() {  
    return repository.findAll();  
}
```

```
public Optional<Student> getStudentById(Long id) {  
    return repository.findById(id);  
}
```

# Implementing Update Operation

## ❖ Purpose:

- Modify an existing entity in the database.

## ❖ Explanation:

- Fetch entity, update fields, and save.

```
public Student updateStudent(Long id, Student updatedStudent) {  
    Student student = repository.findById(id).orElseThrow();  
    student.setName(updatedStudent.getName());  
    student.setEmail(updatedStudent.getEmail());  
    return repository.save(student);  
}
```

# Implementing Delete Operation

## ❖ Purpose:

- Remove an entity from the database.

## ❖ Explanation:

- deleteById() removes the entity with the specified ID.
- Throws exception if ID does not exist.

```
public void deleteStudent(Long id) {  
    repository.deleteById(id);  
}
```

# Creating a REST Controller

## ❖ Purpose:

- Expose CRUD operations via RESTful APIs.

```
@RestController
@RequestMapping("/api/students")
public class StudentController {
    @Autowired
    private StudentService service;

    @PostMapping
    public Student create(@RequestBody Student student) {
        return service.createStudent(student);
    }

    @GetMapping
    public List<Student> getAll() {
        return service.getAllStudents();
    }
}
```

# Full CRUD Example

## ❖ Scenario:

- Manage student records (create, read, update, delete).

```
@RestController
@RequestMapping("/api/students")
public class StudentController {

    @Autowired
    private StudentService service;

    @PostMapping
    public Student create(@RequestBody Student student) {
        return service.createStudent(student);
    }

    @GetMapping("/{id}")
    public Student getById(@PathVariable Long id) {
        return service.getStudentById(id).orElseThrow();
    }

    @PutMapping("/{id}")
    public Student update(@PathVariable Long id, @RequestBody Student student) {
        return service.updateStudent(id, student);
    }

    @DeleteMapping("/{id}")
    public void delete(@PathVariable Long id) {
        service.deleteStudent(id);
    }
}
```

# Testing APIs with Postman

## ❖ Steps:

- Start Spring Boot application.
- Use Postman to send HTTP requests (POST, GET, PUT, DELETE).

## ❖ Example:

- POST: `http://localhost:8080/api/students` with JSON body:

```
{"name": "John Doe", "email": "john@example.com"}
```

# Validation with JPA

## ❖ Purpose:

- Ensure valid data before saving to database.

## ❖ Explanation:

- Use annotations like @NotNull, @Email from javax.validation.

```
public class Student {  
    @NotNull  
    private String name;  
  
    @Email  
    private String email;  
  
}
```



# Introduction to JPA Relationship Annotations

## ❖ Purpose:

- Define relationships between entities (e.g., Student, Department) in JPA.
- Annotations: @OneToOne, @OneToMany, @ManyToOne, @ManyToMany.

## ❖ Key Concepts:

- Owning side: Defines the relationship (owns the foreign key).
- Inverse side: References the owning side (uses mappedBy).
- Cascade and fetch strategies control behavior and performance.

## ❖ Why Important?:

- Enables modeling of complex data relationships in the database.
- Simplifies querying and data management.

# Configuring @OneToOne Relationship

## ❖ Definition:

- One entity instance is related to exactly one instance of another entity.
- Example: A Student has one Address.

## ❖ Explanation:

- @JoinColumn: Specifies the foreign key column (address\_id) in the Student table.
- cascade: Propagates operations (e.g., save, delete) to the related entity.

```
@Entity
public class Student {
    @Id
    @GeneratedValue
    private Long id;
    private String name;

    @OneToOne(cascade = CascadeType.ALL)
    @JoinColumn(name = "address_id")
    private Address address;
}

@Entity
public class Address {
    @Id
    @GeneratedValue
    private Long id;
    private String street;
}
```

# Configuring @OneToMany and @ManyToOne

## ❖ Definition:

- @OneToMany: One entity relates to multiple instances of another (e.g., one Department has many Students).
- @ManyToOne: Many instances relate to one instance (e.g., many Students belong to one Department).

## ❖ Explanation:

- @ManyToOne (owning side): Defines the foreign key (department\_id) in the Student table.
- @OneToMany (inverse side): Uses mappedBy to reference the owning side.

```
@Entity
public class Department {
    @Id
    @GeneratedValue
    private Long id;
    private String name;

    @OneToMany(mappedBy = "department", cascade = CascadeType.ALL)
    private List<Student> students = new ArrayList<>();
}

@Entity
public class Student {
    @Id
    @GeneratedValue
    private Long id;
    private String name;

    @ManyToOne
    @JoinColumn(name = "department_id")
    private Department department;
}
```

# Configuring @ManyToMany Relationship

## ❖ Definition:

- Multiple instances of one entity relate to multiple instances of another.
- Example: Students enroll in multiple Courses, and Courses have multiple Students.

## ❖ Explanation:

- @JoinTable: Defines the join table (student\_course) with foreign keys.
- mappedBy: Specifies the owning side (Student) to avoid duplicate mappings.

```
@Entity
public class Student {
    @Id
    @GeneratedValue
    private Long id;
    private String name;

    @ManyToMany(cascade = CascadeType.ALL)
    @JoinTable(
        name = "student_course",
        joinColumns = @JoinColumn(name = "student_id"),
        inverseJoinColumns = @JoinColumn(name = "course_id")
    )
    private List<Course> courses = new ArrayList<>();
}

@Entity
public class Course {
    @Id
    @GeneratedValue
    private Long id;
    private String title;

    @ManyToMany(mappedBy = "courses")
    private List<Student> students = new ArrayList<>();
}
```

# Cascade and Fetch Strategies in Relationships

## ❖ Cascade:

- Controls propagation of operations (e.g., save, delete) to related entities.
- Options: CascadeType.ALL, PERSIST, MERGE, REMOVE, etc.
- Example: cascade = CascadeType.ALL saves related entities automatically.

## ❖ Fetch Strategies:

- FetchType.LAZY: Loads related data only when accessed (default for @OneToMany, @ManyToMany).
- FetchType.EAGER: Loads related data immediately (default for @ManyToOne, @OneToOne).

```
@OneToMany(mappedBy = "department", cascade = CascadeType.PERSIST, fetch = FetchType.LAZY)  
private List<Student> students;
```

# Introduction to Joins in JPA

## ❖ What are Joins in JPA?:

- Joins in JPA are used to combine data from multiple entities based on relationships.
- Defined in JPQL (Java Persistence Query Language) or Criteria API.
- Support for INNER JOIN, LEFT JOIN, RIGHT JOIN, and implicit joins.

## ❖ Why Use Joins?:

- Retrieve related data in a single query, avoiding multiple database calls.
- Essential for querying associations like @OneToMany or @ManyToMany.

# Inner Join in JPA

## ❖ Definition:

- Returns records that have matching values in both entities.
- Equivalent to SQL INNER JOIN.

## ❖ Explanation:

- Joins Student and Department entities on their relationship.
- Only includes students with a matching department.

```
@Query("SELECT s FROM Student s JOIN s.department d WHERE d.name = :deptName")  
List<Student> findStudentsByDepartment(@Param("deptName") String deptName);
```

# Left Outer Join in JPA

## ❖ Definition:

- Returns all records from the left entity and matching records from the right.
- Non-matching right records are null.

## ❖ Explanation:

- Includes all students, even those without a department.
- Useful for optional relationships.

```
@Query("SELECT s FROM Student s LEFT JOIN s.department d WHERE d.name = :deptName OR d IS NULL")  
List<Student> findStudentsWithOptionalDepartment(@Param("deptName") String deptName);
```



# Right Outer Join in JPA

## ❖ Definition:

- Returns all records from the right entity and matching records from the left.
- Non-matching left records are null.

## ❖ Explanation:

- Includes all departments, even those without students.
- Less common than LEFT JOIN but symmetric.

```
@Query("SELECT d FROM Department d RIGHT JOIN d.students s WHERE s.name = :studentName")  
List<Department> findDepartmentsByStudent(@Param("studentName") String studentName);
```

# Right Outer Join in JPA

## ❖ Definition:

- Returns all records from the right entity and matching records from the left.
- Non-matching left records are null.

```
@Query("SELECT d FROM Department d RIGHT JOIN d.students s WHERE s.name = :studentName")  
List<Department> findDepartmentsByStudent(@Param("studentName") String studentName);
```

# Fetch Joins in JPA

## ❖ Definition:

- Eagerly fetches related entities in a single query to avoid N+1 problem.
- Uses FETCH keyword in JPQL.

## ❖ Explanation:

- Loads departments immediately, preventing lazy loading exceptions.
- Improves performance for read operations.

```
@Query("SELECT s FROM Student s JOIN FETCH s.department d")  
List<Student> findAllStudentsWithDepartments();
```

# Introduction to Paging in JPA

## ❖ Introduction to Paging in JPA

### ❖ What is Paging?:

- Paging allows retrieving large datasets in smaller chunks (pages) to improve performance and usability.
- Essential for applications with large databases to avoid loading all data at once.

### ❖ Key Components in Spring Data JPA:

- Pageable: Interface for pagination and sorting information.
- Page: Represents a page of data with metadata (total pages, total elements).
- Slice: Similar to Page but without total count (faster for large datasets).

### ❖ Why Use Paging?:

- Reduces memory usage, improves response times, and enables features like infinite scrolling.

# Using Pageable in JpaRepository

## ❖ Pageable Interface:

- Created using `PageRequest.of(pageNumber, pageSize, sort)` to specify page index, size, and sorting.

## ❖ Repository Methods:

- Extend `JpaRepository` and add methods returning `Page<T>` or `Slice<T>`.

```
public interface StudentRepository extends JpaRepository<Student, Long> {  
    Page<Student> findAll(Pageable pageable);  
}
```

```
Pageable pageable = PageRequest.of(0, 10, Sort.by("name").ascending());  
Page<Student> studentsPage = repository.findAll(pageable);
```

# Working with Page and Slice

## ❖ Page:

- Provides full pagination info: content, total pages, total elements.
- Example: `studentsPage.getTotalElements()`, `studentsPage.getTotalPages()`.

## ❖ Slice:

- Lighter than Page; no total count (avoids expensive COUNT queries).
- Use for "load more" features where total is not needed.

```
Slice<Student> studentsSlice = repository.findAll(Pageable pageable);  
List<Student> content = studentsSlice.getContent();  
boolean hasNext = studentsSlice.hasNext();
```

# Custom Queries with Paging

## ❖ Custom JPQL Queries:

- Use @Query with Pageable for custom pagination.

## ❖ Explanation:

- Pageable is appended as the last parameter in custom queries.
- Supports sorting and pagination on derived or custom queries.

```
public interface StudentRepository extends JpaRepository<Student, Long> {  
    @Query("SELECT s FROM Student s WHERE s.name LIKE %:name%")  
    Page<Student> findByNameContaining(@Param("name") String name, Pageable pageable);  
}
```

```
Page<Student> results = repository.findByNameContaining("John", pageable);
```

# Best Practices for Paging in JPA

## ❖ Key Practices:

- Use Slice for large datasets to avoid slow COUNT queries.
- Combine with sorting: `Sort.by("field").ascending()` for user-friendly results.
- Handle edge cases: Empty pages, invalid page numbers.
- Use in REST APIs: Return Page metadata in responses for client-side pagination.

## ❖ Common Pitfalls:

- N+1 queries: Use Fetch Joins with paging for relationships.
- Performance: Index columns used in sorting/filters.



# Introduction to Inheritance in JPA Entities

## ❖ Introduction to Inheritance in JPA Entities

### ❖ What is Inheritance in JPA?:

- JPA supports inheritance to model hierarchical entity classes, mapping Java OO inheritance to relational databases.
- Allows subclasses to inherit fields and relationships from a superclass.

### ❖ Inheritance Strategies:

- `@Inheritance(strategy = InheritanceType.SINGLE_TABLE)`: All classes in one table with discriminator.
- `@Inheritance(strategy = InheritanceType.TABLE_PER_CLASS)`: Separate table per concrete class.
- `@Inheritance(strategy = InheritanceType.JOINED)`: Separate table for superclass and each subclass (1:1 relationship via shared primary key).

### ❖ Focus on 1:1 Inheritance:

- Refers to JOINED strategy, where subclass tables link 1:1 to superclass table using shared PK.

### ❖ Reference: JPA Inheritance Specification

# JOINED Inheritance Strategy (1:1 Mapping)

## ❖ **Definition:**

- Superclass has its own table; each subclass has a separate table with only subclass-specific fields.
- Subclass tables reference superclass table via shared primary key (1:1 relationship).

## ❖ **Annotations:**

- @Inheritance(strategy = InheritanceType.JOINED) on superclass.
- @PrimaryKeyJoinColumn optional for customizing join column.

## ❖ **Database Structure:**

- Superclass table: Common fields + PK.
- Subclass table: Subclass fields + PK (foreign key to superclass PK).

## ❖ **Explanation:**

- Queries join tables as needed; supports polymorphism (e.g., querying superclass returns mixed subclass instances).

# Example of 1:1 Inheritance in JPA

## ❖ Generated Tables:

- person: id (PK), name.
- student: id (PK/FK to person.id), major.
- teacher: id (PK/FK to person.id), subject.

## ❖ Explanation:

- Inserting a Student creates rows in both person and student tables with same id.
- Query: SELECT p FROM Person p joins tables to fetch mixed Student/Teacher instances.

## ❖ Reference: Thorben Janssen: JPA Joined Strategy

```
@Entity
@Inheritance(strategy = InheritanceType.JOINED)
public abstract class Person {
    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;
    private String name;
    // Getters and setters
}

@Entity
public class Student extends Person {
    private String major;
    // Getters and setters
}

@Entity
public class Teacher extends Person {
    private String subject;
    // Getters and setters
}
```

# Advantages and Disadvantages of 1:1 Inheritance

## ❖ **Advantages:**

- Normalized database: No redundant fields; easy to add new subclasses.
- Supports polymorphism: Queries on superclass return subclass instances.
- Efficient for reads on specific subclasses (no unnecessary joins).

## ❖ **Disadvantages:**

- Performance overhead: Joins required for superclass queries.
- Complex inserts/updates: Multiple tables involved.
- Not suitable for deep hierarchies due to join complexity.

## ❖ **When to Use:**

- When normalization is important and hierarchies are not too deep.

# Best Practices for 1:1 Inheritance in JPA

## ❖ Key Practices:

- Use @DiscriminatorColumn if needed for explicit type discrimination (though optional in JOINED).
- Optimize queries with Fetch Joins to avoid N+1 issues.
- Index join columns for better performance.
- Test polymorphism: Ensure repositories handle superclass queries correctly.

## ❖ Common Pitfalls:

- Overusing joins in deep hierarchies leading to slow queries.
- Forgetting to generate IDs in superclass.

```
@Query("SELECT p FROM Person p JOIN FETCH p WHERE p.id = :id")  
Person findByIdWithFetch(@Param("id") Long id);
```

# Conclusion and Next Steps

## ❖ **Summary:**

- Learned to build a CRUD application with Spring Boot, JPA, and MySQL.
- Covered entities, repositories, REST APIs, and basics JPA features.

## ❖ **Next Steps:**

- Build a full-stack application with a front-end (e.g., React).

## ❖ **References:**

- Spring Boot
- Spring Data JPA
- JPA Specification